

Transformer Notes

Yue Wu

February 2026

Contents

1	Input Embedding implementation	3
2	Positional Encoding	4
3	Layer Norm	8
4	Multi-head Attention	10
5	Position-wise Feed-Forward Network	14
6	Residual Connection with Pre-Layer Normalization	16
7	Encoder Block	17
8	Transformer Encoder	20
9	Decoder Block	22
10	Transformer Decoder	23
11	Projection Layer	25
12	Full Transformer Architecture	25
13	build_transformer	29




Figure 1: UML diagram

1 Input Embedding implementation

```

1 import torch
2 import torch.nn as nn
3 import math
4
5 class InputEmbeddings(nn.Module):
6     def __init__(self, d_model:int, vocabu_size:int):
7         super().__init__()
8         self.d_model = d_model
9         self.vocabu_size = vocabu_size
10        self.embedding = nn.Embedding(vocabu_size, d_model)
11    def forward(self, x):
12        return self.embedding(x) * math.sqrt(self.d_model)

```

Let

$$E \in \mathbb{R}^{V \times d} \quad (\text{embedding matrix, } V = \text{vocabulary size, } d = \text{model dim})$$

and let the input token indices be

$$X \in \{0, \dots, V - 1\}^{B \times S}$$

with B = batch size and S = sequence length.

We can represent the lookup at each position as a one-hot vector:

$$\mathbf{1}(X_{b,s}) \in \{0, 1\}^V \subset \mathbb{R}^V, \quad \text{where } \mathbf{1}(X_{b,s})_i = \begin{cases} 1 & \text{if } i = X_{b,s}, \\ 0 & \text{otherwise.} \end{cases}$$

Then the embedded tensor (before scaling) can be written as:

$$\tilde{Y}_{b,:} = \mathbf{1}(X_{b,s})^\top E \in \mathbb{R}^d.$$

Stacking over batch and sequence gives:

$$\tilde{Y} \in \mathbb{R}^{B \times S \times d}.$$

Finally the module multiplies by \sqrt{d} (to scale the embeddings):

$$Y = \tilde{Y} \cdot \sqrt{d} \in \mathbb{R}^{B \times S \times d}.$$

Equivalently, if we expand the one-hot vectors into a 3-D tensor $O \in \{0, 1\}^{B \times S \times V}$, we can write

$$\tilde{Y}_{b,:} = O_{b,:}^\top E,$$

2 Positional Encoding

The sinusoidal positional encoding is defined as

$$PE_{(pos, 2i)} = \sin\left(\frac{pos}{10000^{\frac{2i}{d_{model}}}}\right),$$
$$PE_{(pos, 2i+1)} = \cos\left(\frac{pos}{10000^{\frac{2i}{d_{model}}}}\right).$$

Instead of explicitly dividing by $10000^{\frac{2i}{d_{model}}}$, the implementation precomputes:

$$\frac{1}{10000^{\frac{2i}{d_{model}}}}$$

so that the computation can be written as a multiplication:

$$\text{position} \times \text{div_term}.$$

This is numerically stable and computationally efficient.

```
class PositionalEncoding(nn.Module):
    def __init__(self, d_model: int, seq_len: int, dropout: float) -> None:
        super().__init__()
        self.d_model = d_model
        self.seq_len = seq_len
        self.dropout = nn.Dropout(dropout)

        # initialize matrix of shape (seq_len, d_model)
        pe = torch.zeros(seq_len, d_model)

        # create position indices
        position = torch.arange(0, seq_len, dtype=torch.float).unsqueeze(1)

        # numerically stable dividing term
        div_term = torch.exp(
            torch.arange(0, d_model, 2).float() *
            (-math.log(10000.0) / d_model)
        )

        pe[:, 0::2] = torch.sin(position * div_term)
        pe[:, 1::2] = torch.cos(position * div_term)

        # add batch dimension
        pe = pe.unsqueeze(0)
```

```

# register buffer (not updated during backprop)
self.register_buffer('pe', pe)

def forward(self, x):
    x = x + self.pe[:, :x.size(1), :]
    return self.dropout(x)

```

Step 1: Index Vector The expression

$$\text{torch.arange}(0, d_{\text{model}}, 2)$$

produces the vector

$$[0, 2, 4, 6, \dots, d_{\text{model}} - 2].$$

This can be written as

$$2i, \quad i = 0, 1, 2, \dots, \frac{d_{\text{model}}}{2} - 1.$$

Step 2: Multiply by Constant Each element is multiplied by

$$-\frac{\log(10000)}{d_{\text{model}}}.$$

Thus the resulting expression becomes

$$2i \cdot \left(-\frac{\log(10000)}{d_{\text{model}}} \right).$$

Step 3: Apply the Exponential The exponential function is then applied:

$$\text{div_term}_i = \exp \left(2i \cdot \left(-\frac{\log(10000)}{d_{\text{model}}} \right) \right).$$

Step 4: Simplification Using the identity

$$e^{a \log b} = b^a,$$

we obtain

$$\text{div_term}_i = 10000^{-\frac{2i}{d_{\text{model}}}} = \frac{1}{10000^{\frac{2i}{d_{\text{model}}}}}.$$

Final Result Therefore,

$$\text{div_term}_i = \frac{1}{10000^{\frac{2i}{d_{\text{model}}}}}$$

Interpretation The resulting values correspond to geometrically increasing wavelengths:

- Small values of i produce high-frequency oscillations.
- Large values of i produce low-frequency oscillations.
- Each embedding dimension encodes position at a different scale.

Below we explain each line of the implementation in detail.

Initialization

- `super().__init__()`

Initializes the parent `nn.Module` class so that the module properly registers parameters and buffers.

- `self.d_model = d_model`

Stores the embedding dimension.

- `self.seq_len = seq_len`

Stores the maximum sequence length for which positional encodings will be created.

- `self.dropout = nn.Dropout(dropout)`

Creates a dropout layer that will be applied after adding positional encodings.

Creating the Positional Encoding Matrix

- `pe = torch.zeros(seq_len, d_model)`

Creates a tensor of zeros with shape

$$(seq_len, d_model).$$

Each row will correspond to one position in the sequence.

- `position = torch.arange(0, seq_len).unsqueeze(1)`

Creates a column vector of shape

$$(seq_len, 1),$$

containing the position indices:

$$0, 1, 2, \dots, seq_len - 1.$$

- `div_term = torch.exp(
 torch.arange(0, d_model, 2).float() *
 (-math.log(10000.0) / d_model)
)`

Creates a vector containing the scaling factors for each pair of dimensions.

- `torch.arange(0, d_model, 2)` selects even indices.
- The exponential term ensures geometrically increasing wavelengths.
- The division by d_{model} stabilizes the scaling.

The resulting shape is:

$$(d_{\text{model}}/2).$$

Applying Sine and Cosine

- `pe[:, 0::2] = torch.sin(position * div_term)`
Applies the sine function to all even dimensions.
- `pe[:, 1::2] = torch.cos(position * div_term)`
Applies the cosine function to all odd dimensions.

After these operations,

$$pe \in \mathbb{R}^{seq_len \times d_{\text{model}}}.$$

Adding Batch Dimension

- `pe = pe.unsqueeze(0)`

Adds a batch dimension at the front.

New shape:

$$(1, seq_len, d_{\text{model}}).$$

This allows broadcasting across batches.

Registering as Buffer

- `self.register_buffer('pe', pe)`

Registers `pe` as a buffer.

A buffer:

- is saved with the model,
- moves to GPU with the model,
- is not updated during backpropagation.

Forward Pass

- `x = x + self.pe[:, :x.size(1), :]`

Adds positional encodings to the input embeddings.

- `x` has shape (B, S, d_{model}) .
- `self.pe` has shape $(1, seq_len, d_{\text{model}})$.
- Broadcasting expands the first dimension automatically.

- `return self.dropout(x)`

Applies dropout and returns the result.

3 Layer Norm

```
class LayerNormalization(nn.Module):
    def __init__(self, parameters_shape, eps=1e-5):
        super().__init__()
        self.parameters_shape = parameters_shape
        self.eps = eps
        self.gamma = nn.Parameter(torch.ones(parameters_shape))
        self.beta = nn.Parameter(torch.zeros(parameters_shape))

    def forward(self, x):
        dims = [-(i+1) for i in range(len(self.parameters_shape))]

        mean = x.mean(dim=dims, keepdim=True)
        print(f"Mean \n ({mean.size()}) : \n {mean}")

        var = ((x - mean) ** 2).mean(dim=dims, keepdim=True)
        std = (var + self.eps).sqrt()
        print(f"Standard Deviation \n ({std.size()}) : \n {std}")

        y = (x - mean) / std
        print(f"y \n ({y.size()})= \n {y}")

        out = self.gamma * y + self.beta
        return out
```

Layer Normalization normalizes across the feature dimension for each individual sample.
Let

$$x \in \mathbb{R}^{B \times S \times d_{\text{model}}},$$

where:

- B is the batch size,
- S is the sequence length,
- d_{model} is the embedding dimension.

Layer Normalization is applied independently at each position (b, s) across the feature dimension.

Step 1: Compute Mean

For each token vector $x_{b,s,:} \in \mathbb{R}^{d_{\text{model}}}$:

$$\mu_{b,s} = \frac{1}{d_{\text{model}}} \sum_{j=1}^{d_{\text{model}}} x_{b,s,j}.$$

Step 2: Compute Variance

$$\sigma_{b,s}^2 = \frac{1}{d_{\text{model}}} \sum_{j=1}^{d_{\text{model}}} (x_{b,s,j} - \mu_{b,s})^2.$$

Step 3: Normalize

$$\hat{x}_{b,s,j} = \frac{x_{b,s,j} - \mu_{b,s}}{\sqrt{\sigma_{b,s}^2 + \varepsilon}}.$$

Step 4: Scale and Shift

$$y_{b,s,j} = \gamma_j \hat{x}_{b,s,j} + \beta_j$$

Thus, the compact form of LayerNorm is:

$$\text{LayerNorm}(x) = \gamma \odot \frac{x - \mu}{\sqrt{\sigma^2 + \varepsilon}} + \beta$$

where:

- $\gamma \in \mathbb{R}^{d_{\text{model}}}$ is a learnable scaling parameter,
- $\beta \in \mathbb{R}^{d_{\text{model}}}$ is a learnable shift parameter,
- ε is a small constant for numerical stability.

1. The ε Term

$$\varepsilon > 0$$

is added inside the square root to prevent division by zero and ensure numerical stability when the variance is very small.

2. The Scale Parameter γ

After normalization:

$$\mathbb{E}[\hat{x}] = 0, \quad \text{Var}[\hat{x}] = 1.$$

The parameter γ allows the model to rescale features, restoring flexibility and allowing different dimensions to have different learned magnitudes.

3. The Shift Parameter β

Normalization forces zero mean. The parameter β allows the model to shift the features to a non-zero mean if beneficial.

4 Multi-head Attention




Figure 2: Multihead Self Attention

```
class MultiheadAttention(nn.Module):
    def __init__(self, d_model, n_heads, dropout) -> None:
        super().__init__()
```

```

    self.d_model = d_model
    self.n_heads = n_heads

    assert d_model % n_heads == 0, \
        "Model dimension is not divisible by number of heads"

    self.d_k = d_model // n_heads

    self.w_q = nn.Linear(d_model, d_model)
    self.w_k = nn.Linear(d_model, d_model)
    self.w_v = nn.Linear(d_model, d_model)
    self.w_o = nn.Linear(d_model, d_model)

    self.dropout = nn.Dropout(dropout)

@staticmethod
def attention(query, key, value, mask, dropout: nn.Dropout):
    d_k = query.shape[-1]

    # Attention scores
    attention_scores = (query @ key.transpose(-2, -1)) \
        / math.sqrt(d_k)

    if mask is not None:
        attention_scores.masked_fill_(mask == 0, -1e11)

    attention_scores = attention_scores.softmax(dim=-1)

    if dropout is not None:
        attention_scores = dropout(attention_scores)

    qkv = attention_scores @ value

    return qkv, attention_scores

def forward(self, q, k, v, mask):

    # Linear projections
    query = self.w_q(q)
    key = self.w_k(k)
    value = self.w_v(v)

    # Reshape for multi-head attention

```

```

# (B, S, d_model) ->
# (B, S, n_heads, d_k) ->
# (B, n_heads, S, d_k)

query = query.view(query.shape[0], query.shape[1],
                    self.n_heads, self.d_k).transpose(1, 2)

key = key.view(key.shape[0], key.shape[1],
               self.n_heads, self.d_k).transpose(1, 2)

value = value.view(value.shape[0], value.shape[1],
                     self.n_heads, self.d_k).transpose(1, 2)

x, self.attention_scores = self.attention(
    query, key, value, mask, self.dropout
)

# Combine heads
# (B, n_heads, S, d_k) ->
# (B, S, n_heads, d_k) ->
# (B, S, d_model)

x = x.transpose(1, 2).contiguous().view(
    x.shape[0], -1, self.n_heads * self.d_k
)

return self.w_o(x)

```

After the linear projections, the tensors `query`, `key`, and `value` have shape

$$(B, S, d_{\text{model}})$$

where

- B is the batch size,
- S is the sequence length,
- d_{model} is the embedding dimension.

Goal

Multi-head attention splits the embedding dimension into multiple heads:

$$d_{\text{model}} = n_{\text{heads}} \cdot d_k.$$

Each head operates on a subspace of dimension d_k .

Step 1: Reshape

The instruction

```
query.view(B, S, n_heads, d_k)
```

reshapes the tensor from

$$(B, S, d_{\text{model}})$$

to

$$(B, S, n_{\text{heads}}, d_k).$$

This operation does not change the data values. It only reorganizes the embedding dimension into n_{heads} smaller vectors of size d_k .

Step 2: Transpose

The instruction

```
.transpose(1, 2)
```

swaps dimensions 1 and 2, transforming

$$(B, S, n_{\text{heads}}, d_k)$$

into

$$(B, n_{\text{heads}}, S, d_k)$$

Why This Is Necessary

This rearrangement allows attention to be computed independently for each head.

The attention score computation becomes

$$QK^{\top}$$

with shapes:

$$(B, n_{\text{heads}}, S, d_k) \quad \times \quad (B, n_{\text{heads}}, d_k, S)$$

which produces

$$(B, n_{\text{heads}}, S, S).$$

Thus:

- Each head attends independently.
- Matrix multiplication is efficiently batched.
- Different heads can learn different attention patterns.

5 Position-wise Feed-Forward Network

Implementation

```
class FeedForwardBlock(nn.Module):

    def __init__(self, d_model: int, d_ff: int, dropout: float) -> None:
        super().__init__()
        self.linear1 = nn.Linear(d_model, d_ff)
        self.dropout = nn.Dropout(dropout)
        self.linear2 = nn.Linear(d_ff, d_model)

    def forward(self, x):
        return self.linear2(
            self.dropout(
                torch.relu(
                    self.linear1(x)
                )
            )
        )
```

Mathematical Formulation

Let

$$x \in \mathbb{R}^{B \times S \times d_{\text{model}}}.$$

The feed-forward network is applied independently at each position:

$$\text{FFN}(x) = W_2 \sigma(W_1 x + b_1) + b_2,$$

where:

- $W_1 \in \mathbb{R}^{d_{\text{model}} \times d_{\text{ff}}}$,
- $W_2 \in \mathbb{R}^{d_{\text{ff}} \times d_{\text{model}}}$,
- $\sigma(\cdot)$ is the ReLU activation function.

Dropout is applied after the activation.

Difference Between linear1 and linear2

1. First Linear Layer (linear1)

$$\text{linear1} : \mathbb{R}^{d_{\text{model}}} \rightarrow \mathbb{R}^{d_{ff}}$$

This layer:

- Expands the dimensionality.
- Projects each token embedding into a higher-dimensional space.
- Increases model capacity.

Typically:

$$d_{ff} = 4 \times d_{\text{model}}.$$

This is sometimes called the **expansion layer**.

2. Second Linear Layer (linear2)

$$\text{linear2} : \mathbb{R}^{d_{ff}} \rightarrow \mathbb{R}^{d_{\text{model}}}$$

This layer:

- Reduces dimensionality back to d_{model} .
- Projects back to the original embedding space.
- Ensures compatibility with residual connections.

This is sometimes called the **projection layer**.

Intuition

The feed-forward block performs:

Expand → Non-linearity → Compress.

It allows the Transformer to:

- Learn complex feature transformations.
- Increase expressiveness without increasing sequence interaction.
- Process each token independently (position-wise).

Unlike attention, the feed-forward network does not mix information across sequence positions.

6 Residual Connection with Pre-Layer Normalization Implementation

```
class ResidualConnection(nn.Module):
    def __init__(self, features, dropout) -> None:
        super().__init__()
        self.features = features
        self.dropout = nn.Dropout(dropout)
        self.norm = LayerNormalization(features)

    def forward(self, x, sublayer):
        return x + self.dropout(sublayer(self.norm(x)))
```

Mathematical Formulation

Let

$$x \in \mathbb{R}^{B \times S \times d_{\text{model}}}.$$

This block computes:

$$\boxed{\text{Output}(x) = x + \text{Dropout}(\text{Sublayer}(\text{LayerNorm}(x)))}$$

Step-by-Step Explanation

1. Layer Normalization (Pre-Norm)

$$\tilde{x} = \text{LayerNorm}(x)$$

The input is first normalized across the feature dimension. This stabilizes training and improves gradient flow.

2. Sublayer Transformation

$$z = \text{Sublayer}(\tilde{x})$$

The sublayer may be:

- Multi-Head Attention, or
- Feed-Forward Network.

3. Dropout

$$z' = \text{Dropout}(z)$$

Dropout is applied for regularization.

4. Residual Addition

$$y = x + z'$$

The original input is added back to the transformed output.

Why Use Residual Connections?

Residual connections:

- Improve gradient flow in deep networks.
- Prevent vanishing gradients.
- Allow the model to learn identity mappings easily.
- Make very deep Transformers trainable.

Pre-Norm vs Post-Norm

This implementation uses **Pre-Norm**:

$$x + \text{Sublayer}(\text{LayerNorm}(x))$$

The original Transformer paper used **Post-Norm**:

$$\text{LayerNorm}(x + \text{Sublayer}(x)).$$

Modern Transformers typically use Pre-Norm because it leads to more stable training.

7 Encoder Block




Figure 3: Encoder Block

Implementation

```
class EncoderBlock(nn.Module):
    def __init__(self, self_attention_block: MultiheadAttention,
                 feed_forward_block: FeedForwardBlock,
                 features: int,
                 dropout: float) -> None:
        super().__init__()
        self.self_attention_block = self_attention_block
        self.feed_forward_block = feed_forward_block

        # Two residual connections:
        # one for attention, one for feed-forward
        self.residual_connections = nn.ModuleList(
            [ResidualConnection(features, dropout) for _ in range(2)])
    )

    def forward(self, x, src_mask):

        # Self-attention + residual
        x = self.residual_connections[0](
            x,
            lambda x: self.self_attention_block(x, x, x, src_mask)
        )

        # Feed-forward + residual
        x = self.residual_connections[1](
            x,
            self.feed_forward_block
        )

    return x
```

Mathematical Formulation

Let

$$x \in \mathbb{R}^{B \times S \times d_{\text{model}}}.$$

An encoder block consists of two sublayers:

1. Multi-head self-attention
2. Position-wise feed-forward network

Step 1: Self-Attention with Residual (Pre-Norm)

$$x^{(1)} = x + \text{Dropout}(\text{SelfAttention}(\text{LayerNorm}(x)))$$

Since this is self-attention:

$$Q = K = V = x.$$

Step 2: Feed-Forward with Residual (Pre-Norm)

$$x^{(2)} = x^{(1)} + \text{Dropout}(\text{FFN}(\text{LayerNorm}(x^{(1)})))$$

Final Output

$$\boxed{\text{EncoderBlock}(x) = x^{(2)}}$$

Key Properties

- Shape is preserved:
 $(B, S, d_{\text{model}}) \rightarrow (B, S, d_{\text{model}})$
- Attention mixes information across sequence positions.
- The feed-forward network processes each position independently.
- Residual connections improve gradient flow and stabilize deep training.

8 Transformer Encoder

Implementation

```
class Encoder(nn.Module):  
    def __init__(self, features: int, layers: nn.ModuleList):  
        super().__init__()  
        self.layers = layers  
        self.features = features  
        self.norm = LayerNormalization(features)  
  
    def forward(self, x, mask):  
        for layer in self.layers:  
            x = layer(x, mask)  
        return self.norm(x)
```

Mathematical Formulation

Let

$$x^{(0)} \in \mathbb{R}^{B \times S \times d_{\text{model}}}$$

be the embedded input sequence.

Assume the encoder contains N stacked encoder blocks:

$$\{\text{EncoderBlock}_1, \dots, \text{EncoderBlock}_N\}.$$

The encoder applies them sequentially:

$$x^{(1)} = \text{EncoderBlock}_1(x^{(0)}, \text{mask})$$

$$x^{(2)} = \text{EncoderBlock}_2(x^{(1)}, \text{mask})$$

⋮

$$x^{(N)} = \text{EncoderBlock}_N(x^{(N-1)}, \text{mask})$$

After the final block, a Layer Normalization is applied:

$\text{Encoder}(x^{(0)}) = \text{LayerNorm}(x^{(N)})$

Explanation

1. Stacked Structure

The loop

```
for layer in self.layers:  
    x = layer(x, mask)
```

applies N encoder blocks sequentially. Each block refines the representation of the sequence.

2. Shape Preservation

Throughout the encoder:

$$(B, S, d_{\text{model}}) \longrightarrow (B, S, d_{\text{model}})$$

The dimensionality remains constant. Only the representation is transformed.

3. Final Layer Normalization

$$\text{LayerNorm}(x^{(N)})$$

stabilizes the output before it is passed to the decoder or to subsequent components of the Transformer.

9 Decoder Block




Figure 4: Encoder-Decoder

Implementation

```
class DecoderBlock(nn.Module):

    def __init__(self,
                 features: int,
                 self_attention_block: MultiheadAttention,
                 cross_attention_block: MultiheadAttention,
                 feed_forward_block: FeedForwardBlock,
                 dropout: float) -> None:

        super().__init__()
        self.self_attention_block = self_attention_block
        self.cross_attention_block = cross_attention_block
        self.feed_forward_block = feed_forward_block

        # Three residual connections:
        # 1) masked self-attention
        # 2) cross-attention
        # 3) feed-forward
        self.residue_connections = nn.ModuleList(
            [ResidualConnection(features, dropout) for _ in range(3)])
    )
```

```

def forward(self, x, encoder_output, src_mask, tgt_mask):

    # Masked self-attention
    x = self.residue_connections[0](
        x,
        lambda x: self.self_attention_block(x, x, x, tgt_mask)
    )

    # Cross-attention (encoder-decoder attention)
    x = self.residue_connections[1](
        x,
        lambda x: self.cross_attention_block(
            x, encoder_output, encoder_output, src_mask
        )
    )

    # Feed-forward
    x = self.residue_connections[2](
        x,
        self.feed_forward_block
    )

return x

```

10 Transformer Decoder

Implementation

```

class Decoder(nn.Module):
    def __init__(self, layers: nn.ModuleList, features: int):
        super().__init__()
        self.layers = layers
        self.norm = LayerNormalization(features)

    def forward(self, x, encoder_output, tgt_mask, src_mask):
        for layer in self.layers:
            x = layer(x, encoder_output, src_mask, tgt_mask)
        return self.norm(x)

```

Explanation of the Decoder

The Decoder is composed of N stacked DecoderBlocks followed by a final Layer Normalization.

Let

$$x^{(0)} \in \mathbb{R}^{B \times S_{\text{tgt}} \times d_{\text{model}}}$$

be the target input embeddings, and let

$$H^{\text{enc}} \in \mathbb{R}^{B \times S_{\text{src}} \times d_{\text{model}}}$$

be the encoder output.

Stacked Decoder Blocks

The loop

```
for layer in self.layers:
    x = layer(x, encoder_output, src_mask, tgt_mask)
```

applies each DecoderBlock sequentially:

$$x^{(1)} = \text{DecoderBlock}_1(x^{(0)}, H^{\text{enc}})$$

$$x^{(2)} = \text{DecoderBlock}_2(x^{(1)}, H^{\text{enc}})$$

⋮

$$x^{(N)} = \text{DecoderBlock}_N(x^{(N-1)}, H^{\text{enc}})$$

Each DecoderBlock contains:

- Masked self-attention (prevents attending to future tokens),
- Cross-attention (queries the encoder output),
- Feed-forward network,
- Residual connections and Layer Normalization.

Final Layer Normalization

After the final block, a LayerNorm is applied:

$\text{Decoder}(x^{(0)}) = \text{LayerNorm}(x^{(N)})$

This normalization stabilizes the final representation before it is passed to the output projection layer.

Shape Preservation

Throughout the decoder, the tensor shape remains constant:

$$(B, S_{\text{tgt}}, d_{\text{model}}) \longrightarrow (B, S_{\text{tgt}}, d_{\text{model}})$$

Only the internal representation is refined at each layer.

Role of Masks

- The **target mask** ensures autoregressive behavior by preventing attention to future tokens.
- The **source mask** ensures padding tokens in the encoder output are ignored during cross-attention.

11 Projection Layer

Implementation

```
class ProjectionLayer(nn.Module):  
  
    def __init__(self, d_model, vocab_size):  
        super().__init__()  
        self.d_model = d_model  
        self.vocab_size = vocab_size  
        self.proj = nn.Linear(self.d_model, self.vocab_size)  
  
    def forward(self, x):  
        # (batch, seq_len, d_model) --> (batch, seq_len, vocab_size)  
        return self.proj(x)
```

12 Full Transformer Architecture

Implementation

```
class Transformer(nn.Module):  
  
    def __init__(self,  
                 encoder: Encoder,  
                 decoder: Decoder,  
                 src_embed: InputEmbeddings,  
                 tgt_embed: InputEmbeddings,  
                 src_pos: PositionalEncoding,
```

```

tgt_pos: PositionalEncoding,
projection_layer: ProjectionLayer) -> None:

super().__init__()
self.encoder = encoder
self.decoder = decoder
self.src_embed = src_embed
self.tgt_embed = tgt_embed
self.src_pos = src_pos
self.tgt_pos = tgt_pos
self.projection_layer = projection_layer

def encode(self, src, src_mask):
    src = self.src_embed(src)
    src = self.src_pos(src)
    return self.encoder(src, src_mask)

def decode(self, tgt, encoder_output, src_mask, tgt_mask):
    tgt = self.tgt_embed(tgt)
    tgt = self.tgt_pos(tgt)
    return self.decoder(tgt, encoder_output, src_mask, tgt_mask)

def project(self, x):
    return self.projection_layer(x)

```

The Transformer model is composed of three main components:

1. Source embedding and encoder
2. Target embedding and decoder
3. Final linear projection to vocabulary space




Figure 5: Encoder-Decoder

1. Source Encoding

Let the source input be

$$\text{src} \in \mathbb{R}^{B \times S_{\text{src}}},$$

where:

- B is the batch size,
- S_{src} is the source sequence length.

Embedding Each token index is mapped to a vector of dimension d_{model} :

$$E_{\text{src}} = \text{Embed}_{\text{src}}(\text{src}) \in \mathbb{R}^{B \times S_{\text{src}} \times d_{\text{model}}}.$$

Positional Encoding Positional information is added:

$$\tilde{E}_{\text{src}} = E_{\text{src}} + PE_{\text{src}}.$$

Encoder The encoder transforms the representation:

$$H^{\text{enc}} = \text{Encoder}(\tilde{E}_{\text{src}}) \in \mathbb{R}^{B \times S_{\text{src}} \times d_{\text{model}}}.$$

The encoder output contains contextualized representations of the source sequence.

2. Target Decoding

Let the target input be

$$\text{tgt} \in \mathbb{R}^{B \times S_{\text{tgt}}}.$$

Target Embedding

$$E_{\text{tgt}} = \text{Embed}_{\text{tgt}}(\text{tgt}) \in \mathbb{R}^{B \times S_{\text{tgt}} \times d_{\text{model}}}.$$

Add Positional Encoding

$$\tilde{E}_{\text{tgt}} = E_{\text{tgt}} + PE_{\text{tgt}}.$$

Decoder The decoder uses both the target representation and the encoder output:

$$H^{dec} = \text{Decoder}(\tilde{E}_{\text{tgt}}, H^{enc}).$$

The decoder applies:

- Masked self-attention,
- Cross-attention with encoder outputs,
- Feed-forward transformations.

The resulting tensor has shape

$$H^{dec} \in \mathbb{R}^{B \times S_{\text{tgt}} \times d_{\text{model}}}.$$

3. Projection to Vocabulary

The final step maps the decoder output to vocabulary logits:

$\text{logits} = H^{dec}W + b$

where

$$W \in \mathbb{R}^{d_{\text{model}} \times V}, \quad b \in \mathbb{R}^V,$$

and V is the vocabulary size.

The final output shape is

$$(B, S_{\text{tgt}}, V).$$

Applying softmax gives probabilities over the vocabulary:

$$\text{Softmax}(\text{logits}).$$

Overall Transformer Function

The entire Transformer can be written compactly as:

$$\text{Transformer}(\text{src}, \text{tgt}) = \text{Projection}\left(\text{Decoder}(\text{Encoder}(\text{src}), \text{tgt})\right)$$

Key Properties

- The encoder processes the source sequence.
- The decoder generates the target sequence autoregressively.
- Attention enables global interaction across tokens.
- Dimensionality d_{model} remains constant throughout the model.

13 build_transformer

Implementation

```
def build_transformer(src_vocab_size: int,
                      tgt_vocab_size: int,
                      src_seq_len: int,
                      tgt_seq_len: int,
                      d_model: int = 512,
                      N: int = 6,
                      h: int = 8,
                      dropout: float = 0.1,
                      d_ff: int = 2048):

    # Create the embedding layers
    src_embed = InputEmbeddings(d_model, src_vocab_size)
    tgt_embed = InputEmbeddings(d_model, tgt_vocab_size)

    # Create the positional encoding layers
    src_pos = PositionalEncoding(d_model, src_seq_len, dropout)
    tgt_pos = PositionalEncoding(d_model, tgt_seq_len, dropout)

    # Create the encoder blocks
    encoder_blocks = []
    for _ in range(N):
        encoder_self_attention_block = MultiheadAttention(d_model, h, dropout)
        feed_forward_block = FeedForwardBlock(d_model, d_ff, dropout)
        encoder_block = EncoderBlock(encoder_self_attention_block,
```

```

                feed_forward_block,
                d_model,
                dropout)
encoder_blocks.append(encoder_block)

# Create the decoder blocks
decoder_blocks = []
for _ in range(N):
    decoder_self_attention_block = MultiheadAttention(d_model, h, dropout)
    decoder_cross_attention_block = MultiheadAttention(d_model, h, dropout)
    feed_forward_block = FeedForwardBlock(d_model, d_ff, dropout)
    decoder_block = DecoderBlock(d_model,
                                 decoder_self_attention_block,
                                 decoder_cross_attention_block,
                                 feed_forward_block,
                                 dropout)
    decoder_blocks.append(decoder_block)

# Create the encoder and decoder
encoder = Encoder(d_model, nn.ModuleList(encoder_blocks))
decoder = Decoder(nn.ModuleList(decoder_blocks), d_model)

# Create the projection layer
projection_layer = ProjectionLayer(d_model, tgt_vocab_size)

# Create the transformer
transformer = Transformer(encoder, decoder,
                         src_embed, tgt_embed,
                         src_pos, tgt_pos,
                         projection_layer)

# Initialize the parameters (Xavier / Glorot for matrices)
for p in transformer.parameters():
    if p.dim() > 1:
        nn.init.xavier_uniform_(p)

return transformer

```

What this function constructs

- **Input embeddings:** $\text{Embed}_{\text{src}}, \text{Embed}_{\text{tgt}}$ of shape mapping token indices $\rightarrow \mathbb{R}^{d_{\text{model}}}$.
- **Positional encodings:** fixed sinusoidal positional layers for source/target with maximum lengths $\text{src_seq_len}, \text{tgt_seq_len}$.

- **Encoder stack:** N copies of EncoderBlock. Each block contains self-attention and position-wise feed-forward sublayers (with residual + LayerNorm).
- **Decoder stack:** N copies of DecoderBlock. Each block contains masked self-attention, cross-attention and feed-forward (with residual + LayerNorm).
- **Projection:** linear layer $W \in \mathbb{R}^{d_{\text{model}} \times V_{\text{tgt}}}$ mapping decoder outputs to vocabulary logits.
- **Parameter initialization:** Xavier (Glorot) uniform initialization is applied to all parameters with $\text{dim} > 1$ (i.e., weight matrices).

Shapes / notation

- Input tokens: $\text{src} \in \mathbb{Z}^{B \times S_{\text{src}}}$, $\text{tgt} \in \mathbb{Z}^{B \times S_{\text{tgt}}}$.
- After embedding: $\in \mathbb{R}^{B \times S \times d_{\text{model}}}$.
- Encoder / decoder keep the hidden shape (B, S, d_{model}) .
- Final logits: $(B, S_{\text{tgt}}, V_{\text{tgt}})$.